



Voter Analytics Pipeline

Using Simulated Voter Records | [Goodparty.org](https://goodparty.org)

Agenda



Executive Summary



**Business Problem &
Objective**



Technical Approach



**Orchestration &
Lineage**



**Data Quality &
Governance**



**Business Impact &
Sample Analytics**



Surfacing Insights



**What's Next?
Possible Extensions**

Production-Ready Voter Analytics Platform

Scalable Data Ingestion and Transformation to Enable Deep Electoral Analytics



Task

- Build a daily analytics pipeline
- Ingest raw voter records
- Curate analytics datasets for decision-making

Project Scope: Take-home assessment

Timeline: 4-6 hours estimated

Tech Stack: Airflow | dbt | DuckDB | Streamlit

Automated



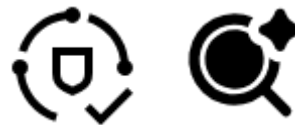
2 Python Processors

3 Airflow Dags

8 dbt Models

1 Streamlit Prototype

Quality



Statically Typed

Unit Tested

dbt Contracted

Auditable

Insights



5 Marts

4 Dims

2 Facts

Built for Analysis

From Voter Records to Campaign Strategy



Rich Electoral Analytics fueled by a pipeline built for growth

Challenges

Periodic voter file
(sub-weekly update)



No guaranteed
validation



Integrity risk
in input files



No historic
election context



Required	Delivered	
Idempotency	✓	Automated EL-processor handling incremental ingestion
Intermediate	✓	Sanitized raw records in contracted DIM and STAGE layers
Mart	✓	Curated tables for core aggregations and targeting
Best Practice	✓	Modular, tested, documented, and reproducible

Extensions (Beyond MVP)

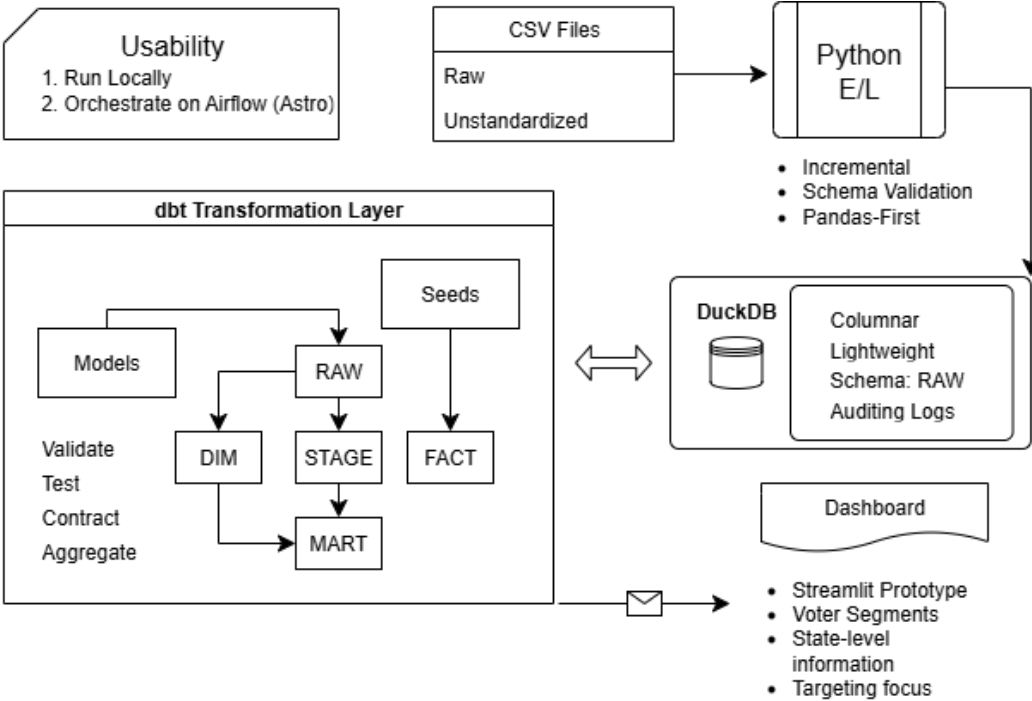
Scalable Build System	3 DAGs Setup > [Daily Pipeline, Monthly Seed]
Integrated Election Calendar	MIT Historic Google Civic API Federal Schedule
Behavioral Segmentation	6 engagement tiers, derived opportunity scoring
Production Patterns	Data contracting, custom macro(s)
Prototype Dashboard	Streamlit app surfacing 8 interactive visualizations

Technical Approach

Pandas ELT | Medallion Architecture



Layer	Technology	Rationale
Orchestration	Airflow (Astro) Cosmos	Rapid local development, dbt integration, portable
Ingestion	Python Pandas	<div>✓ MVP: no-frills basic load strategy</div> <div>🚀 Extension: MD5 deduplication, schema validation, error thresholding, batch processing</div>
Storage	DuckDB	Analytics-optimized, embedded, no infrastructure overhead, good for early phase
Transformation	dbt	<div>✓ MVP: essential dbt validation</div> <div>🚀 Extension: Ingestion unit testing, type-safe interfaces, dbt-expectations tests, macro-ready</div>
Visualization	Streamlit + Plotly	Self-service analytics, no BI tool dependencies, good for prototyping

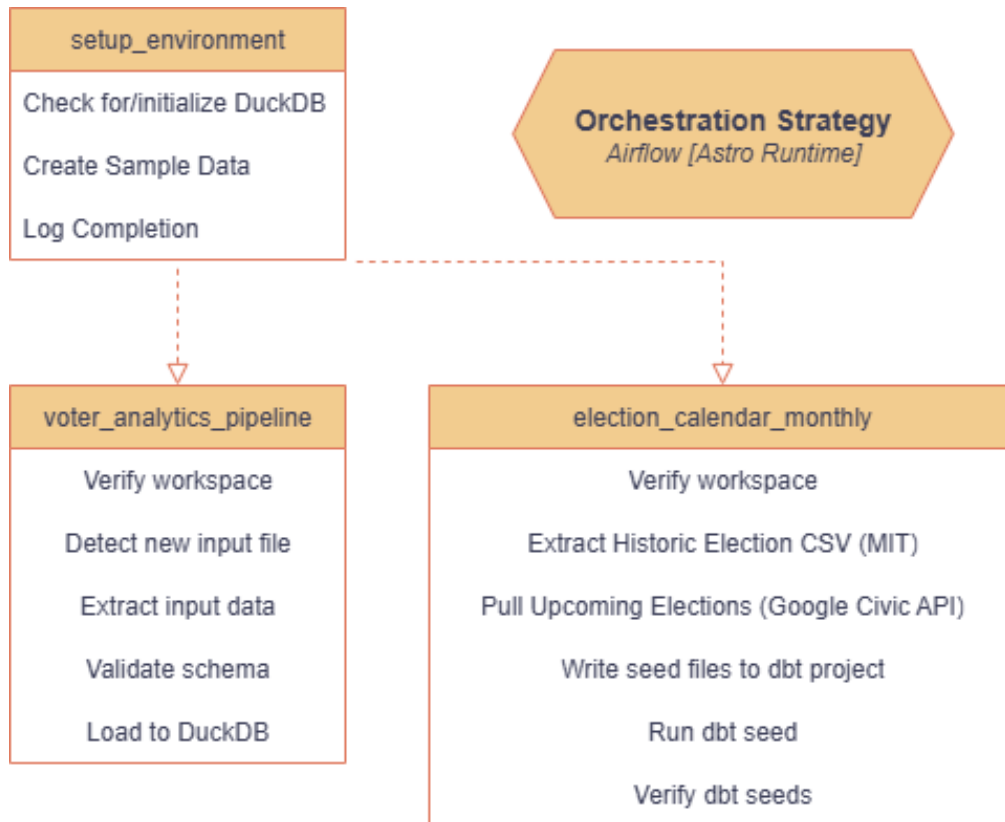


Pipeline Stats:

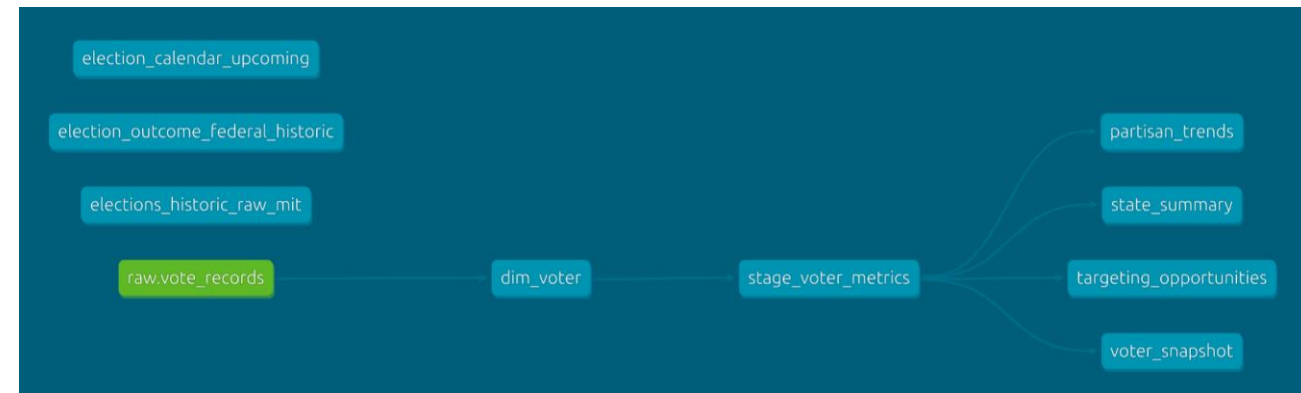
- 3 Airflow DAGs
- 15+ dbt models (dim, stage, mart)
- 15+ data quality tests
- 4 production mart tables
- 2 election seed files (historic + upcoming)

Orchestration & Lineage

Layered approach for flexibility and reusability



dbt Lineage



Ready to leverage new seeds

1. DIM_VOTER: cleans and standardizes inbound vote_records upon landing in RAW
2. STAGE_VOTER_METRICS: prepares voter records for aggregation
3. Marts: partisanship, demographics, regional insight-ready

Data Quality & Governance

Production-Grade Testing, Contracts, Type Safety, and Referential Integrity at All Stages



Schema Enforcement

MVP: Validate key fields

Extension: type-safe EL pattern and dbt contracts

Quality Assurance

MVP: Tests for errors and warnings

Extension: EL unit-testing, dbt column-tests, dbt-expectations

ETL Safeguards

MVP: Handle Incremental Loads

Extensions:

Quick error check |
>5% malformed records

Schema Validation |
Enforces 10 expected fields

Record Deduplication |
MD5 Hash:
[ID, First/Last Name, Email]

Batch Process |
Default 1,000 records

Garbage Collection |
Closes connections
Respects DuckDB 1-thread

```
# Example from voter_snapshot
contract:
  enforced: true
columns:
  - name: total_voters
    data_type: bigint
  - name: pct_current_voters
    data_type: decimal(5,2)
```

Test Coverage	Covered Domains
DB Config + IO	Storage
Pre-ingest Typing	All inbound fields
Uniqueness	Voter IDs, State Codes
Null-checking	IDs, demographics, marts
Regex Validation	Emails, State Codes
Range Validation	Age (18-120), dates, percentages (0-100)
Accepted Values	Parties, States

Business Impact & Sample Analytics

Evolving Raw Records into Data Strategy



PROD_MART.VOTER_SNAPSHOT

Purpose: current voter composition

MVP: voter count by state, party

Extensions:

Behavioral segments | Engagement tiers

- 6 Engagement Segments:
- Current Voter (participated recently)
 - Missed Last Election (lapsed once)
 - Occasional Voter (2-3 lapses)
 - Infrequent (4-6 lapses)
 - Dormant (7+ lapses)
 - Never Voted

Sample Insight:*

"Pennsylvania has 12,500 high-value 'Missed Last Election' target Democrats"

PROD_MART.TARGETING_OPPORTUNITIES

Purpose: ranked segments for GOTV campaigns

MVP: not required

Extensions:

Opportunity score algorithm prioritizing recency

- Opportunity Score (0-100):
- 40% weight: Recent lapsers (1 election)
 - 30% weight: Medium lapsers (2-3)
 - 20% weight: Registration tenure
 - 10% weight: Segment size

Sample Insight:

"Top 20 segments represent 45,000 recoverable voters with 78% recent engagement history"

PROD_MART.PARTISAN_TRENDS

Purpose: time series participation analysis

MVP: not required

Extensions:

Turnout trends over 9 election cycles (2008-2024)

Sample Insight:

"Independent voter participation dropped 18% from 2020 to 2022 midterms suggesting mobilization gap"

PROD_MART.STATE_SUMMARY

Purpose: geographic competitive landscape

MVP: voter count by state

Extensions:

Partisan lean classification |
Engagement Opportunity Scoring

- Partisan Lean Categories:
- Strong Dem/Rep (>10% margin)
 - Lean Dem/Rep (5-10%)
 - Competitive (2-5%)
 - Highly Competitive (<2%)

Sample Insight:

"3 highly competitive states (NC, GA, AZ) have 35% recoverable voter populations"

Strategic Impact

1. Trend participation rates across cycles and partisanship
2. Identify high-priority re-engagement targets among core voter segments
3. Prioritize competitive states for resource allocation
4. Segment voters for tailored messaging

Surfacing Insights

Interactive Streamlit dashboard prototype, no SQL required



🎯 Top Targeting Opportunities

State	Age Group	Party	Opportunity Score
WA	50-64	Republican	<div><div></div></div> 56.7%
OR	30-49	Independent	<div><div></div></div> 56.7%
NH	30-49	Democrat	<div><div></div></div> 56.7%
AK	30-49	Independent	<div><div></div></div> 53.4%
PA	30-49	Republican	<div><div></div></div> 53.4%
OR	50-64	Democrat	<div><div></div></div> 53.4%
NH	50-64	Republican	<div><div></div></div> 53.4%
NJ	30-49	Democrat	<div><div></div></div> 47.5%
DE	30-49	Democrat	<div><div></div></div> 47.5%
CA	30-49	Democrat	<div><div></div></div> 47.5%

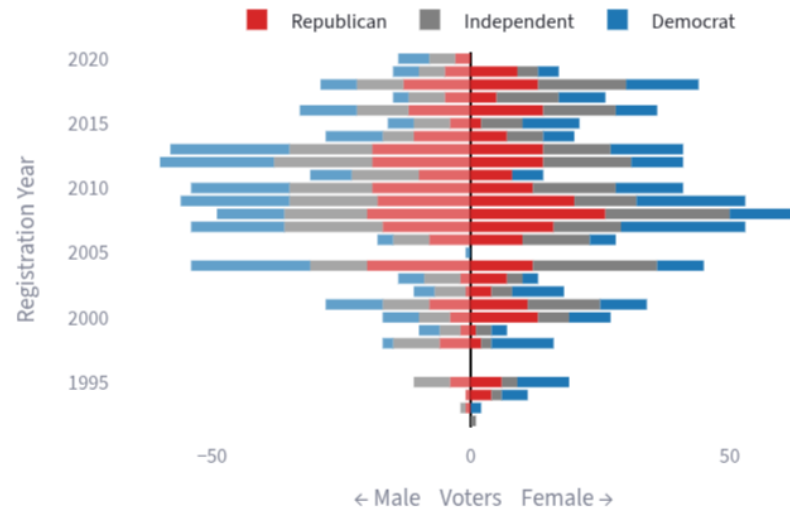
🎨 Dashboard Features:

- 8 interactive charts (bar, line, heatmap, diverging, pie)
- Real-time filtering (state, party, engagement tier)
- Drill-down from state → demographic segments
- Export-ready tables for campaign teams

Voter Registrations by Year

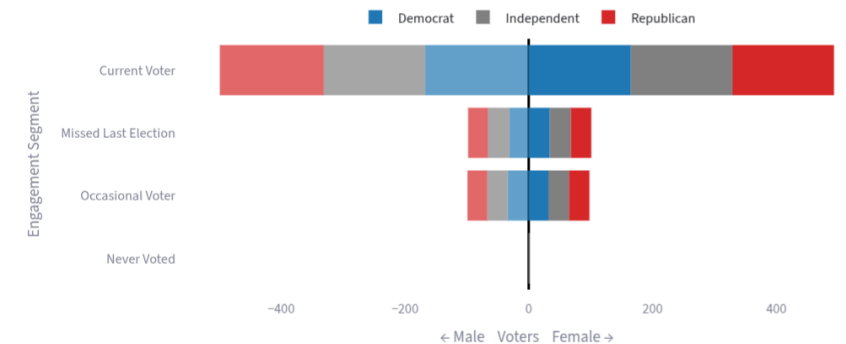


Partisanship by Gender Over Time



🎯 Targeting by Engagement & Demographics

Gender & Engagement by Party



What's Next?

Scaling the platform from MVP to enterprise



Phase 1 [Complete]

- Idempotent ETL pipeline
- Medallion architecture
(raw->dim->stage->mart)
- Quality testing and data contracts
- Behavioral segmentation and opportunity scoring
- Interactive dashboard
- Portable distribution

Potential Extensions

• Phase 2: Production Hardening [2-4 weeks]

- Refactor/integrate into existing platform (DB, BI, docs)
- Quality Improvements: Deeper testing, automated Airflow monitoring (Slack/email), audit suite
- Performance improvements: partitioning, snapshotting (e.g. address changes), query optimizations (clustering, indexes, materialized views)
- Security & Compliance: PII hashing, RBAC, privacy compliance

• Phase 3: Advanced Analytics [2-3 months]

- Predictive modeling: turnout prediction, churn risk modeling
- Enriched dimensions: household clustering, social listening
- Real-Time Operations: automate anomaly detection. CDC from upstream systems

Thank You | Discussion



How does the team balance velocity and technical rigor when making architectural decisions?
Framework for evaluating build-vs-buy and 'good enough' vs. production-hardened tradeoffs?

What are the biggest data quality challenges the team is facing today?
How do you handle schema evolution in production?
How does the team approach net-new data models? (For example, Serve)
When something breaks in production, what is the incident response process?

What strategic or operational decisions are hardest to make with data you have today?
How do you balance data investments between Win (mature) and Serve (emerging)?

Materials

[GitHub Repository](#)

[PEW Research Party Affiliation
Fact Sheet \(NOPRS\)](#)

Socials

[Find me on LinkedIn](#)

[Check out my GitHub](#)

[Read on Medium](#)